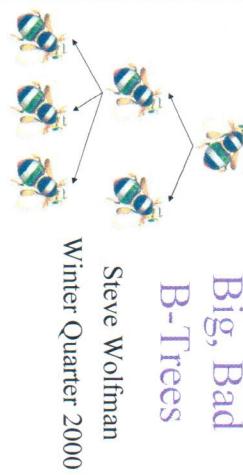


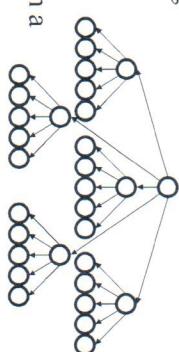
Today's Outline

CSE 326: Data Structures
Lecture #9
Big, Bad
B-Trees



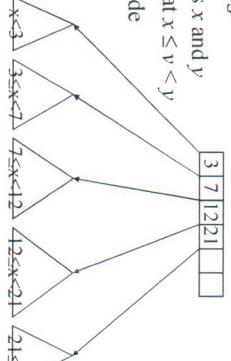
Mary Search Tree

- Maximum branching factor of M



B-Trees

- B-Trees are specialized *Mary* search trees
- Each node has many keys
 - subtree between two keys x and y
 - contains values v such that $x \leq v < y$
 - binary search within a node to find correct subtree
- Each node takes one full $\{page, block, line\}$ of memory



TM 1 2 3

B-Tree Properties^{*}

- Properties
 - maximum branching factor of M
 - **the root has between 2 and M children or at most L keys**
 - other internal nodes have between $\lceil M/2 \rceil$ and M children
 - internal nodes contain only search keys (no data)
 - smallest datum between search keys x and y equals x
 - each (non-root) leaf contains between $\lceil L/2 \rceil$ and L keys
 - all leaves are at the same depth
 - Result
 - tree is $\Theta(\log n)$ deep
 - all operations run in $\Theta(\log n)$ time
 - operations pull in about M or L items at a time
- ^{*}These are technically B⁺-Trees

B-Tree Properties

- Properties
 - maximum branching factor of M
 - the root has between 2 and M children or at most L keys
 - other internal nodes have between $\lceil M/2 \rceil$ and M children
 - internal nodes contain only search keys (no data)
 - smallest datum between search keys x and y equals x
 - each (non-root) leaf contains between $\lceil L/2 \rceil$ and L keys
 - all leaves are at the same depth
- Result
 - tree is $\Theta(\log n)$ deep
 - all operations run in $\Theta(\log n)$ time
 - operations pull in about M or L items at a time

B-Tree Properties

- Properties
 - maximum branching factor of M
 - the root has between 2 and M children or at most L keys
 - other internal nodes have between $\lceil M/2 \rceil$ and M children
 - internal nodes contain only search keys (no data)
 - smallest datum between search keys x and y equals x
 - each (non-root) leaf contains between $\lceil L/2 \rceil$ and L keys
 - all leaves are at the same depth
- Result
 - tree is $\Theta(\log_M n)$ deep
 - all operations run in $\Theta(\log_M n)$ time
 - operations pull in about $M/2$ or $L/2$ items at a time

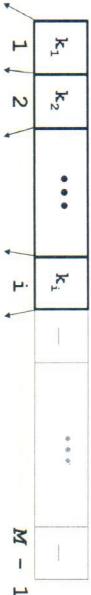
B-Tree Properties

- Properties
 - maximum branching factor of M
 - the root has between 2 and M children or at most L keys
 - other internal nodes have between $\lceil M/2 \rceil$ and M children
 - internal nodes contain only search keys (no data)
 - smallest datum between search keys x and y equals x
 - each (non-root) leaf contains between $\lceil L/2 \rceil$ and L keys
 - all leaves are at the same depth
- Result
 - tree is $\Theta(\log_M n)$ deep
 - all operations run in $\Theta(\log_M n)$ time
 - operations pull in about $M/2$ or $L/2$ items at a time

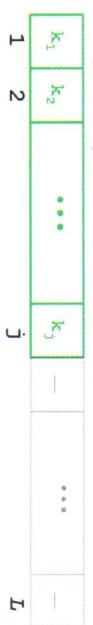
B-Tree Properties

B-Tree Nodes

- Internal node
 - i search keys; $i+1$ subtrees; $M - i - 1$ inactive entries

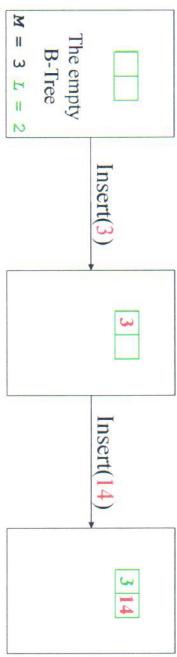


- Leaf
 - j data keys; $L - j$ inactive entries

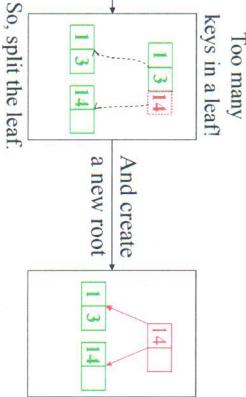


Making a B-Tree

Now, Insert(1)?

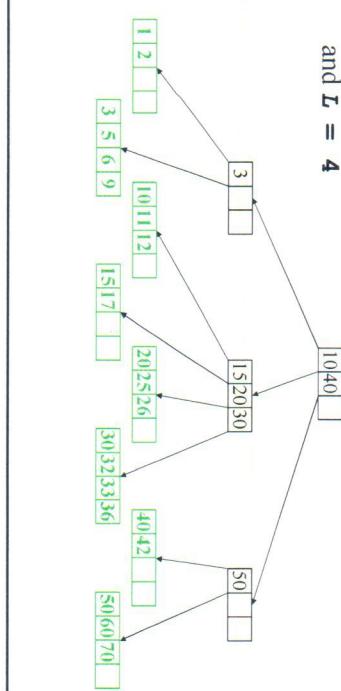


Splitting the Root



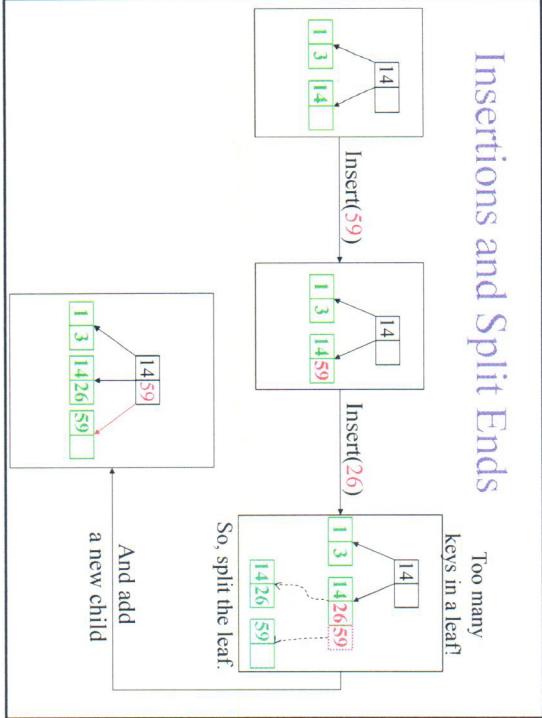
Example

B-Tree with $M = 4$
and $L = 4$



Insertions and Split Ends

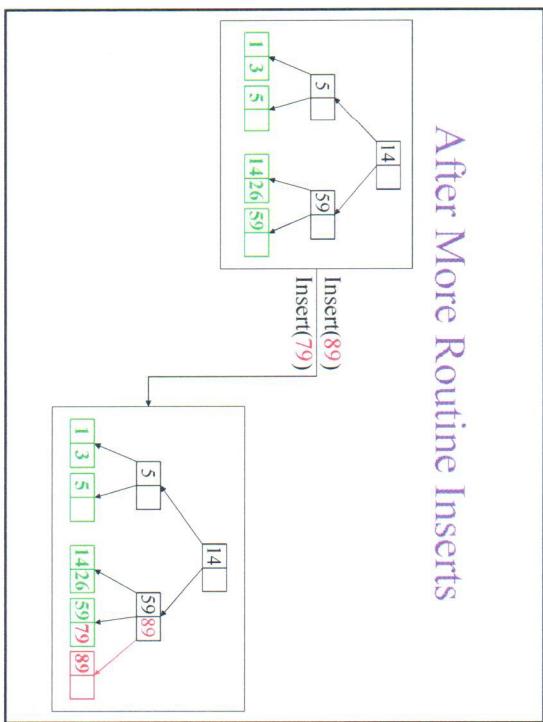
Too many keys in a leaf!



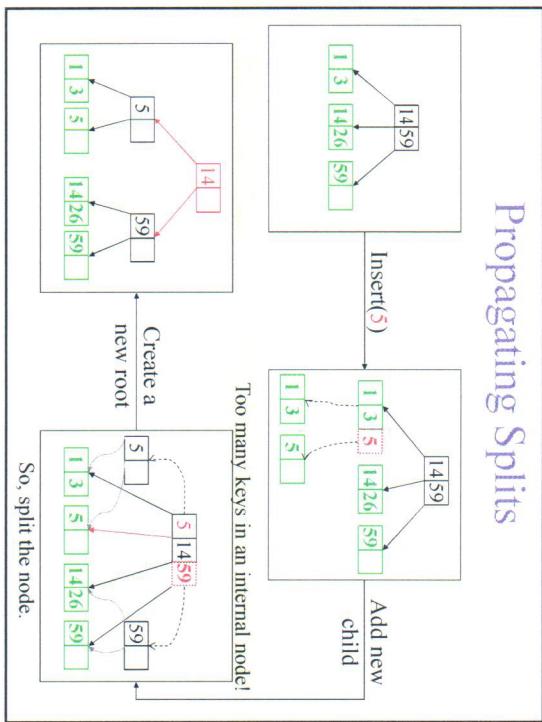
Insertion in Boring Text

- Insert the key in its leaf
 - If the leaf ends up with $L+1$ with $M+1$ items, **overflow!**
 - Split the leaf into two nodes:
 - original with $\lceil (L+1)/2 \rceil$ items
 - new one with $\lfloor (L+1)/2 \rfloor$ items
 - Add the new child to the parent
 - If the parent ends up with $M+1$ items, **overflow!**
 - Split an overflowed root in two and hang the new nodes under a new root
- This makes the tree deeper!

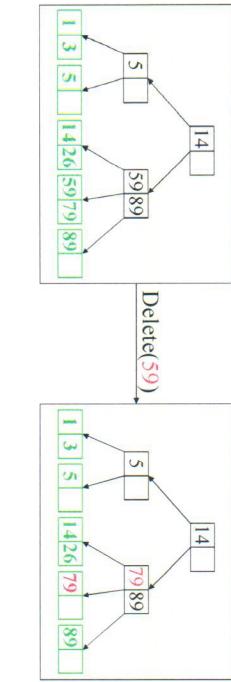
After More Routine Inserts



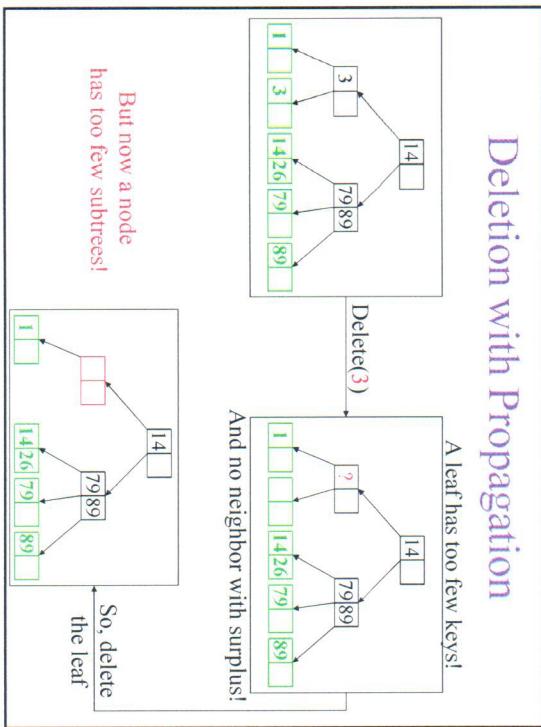
Propagating Splits



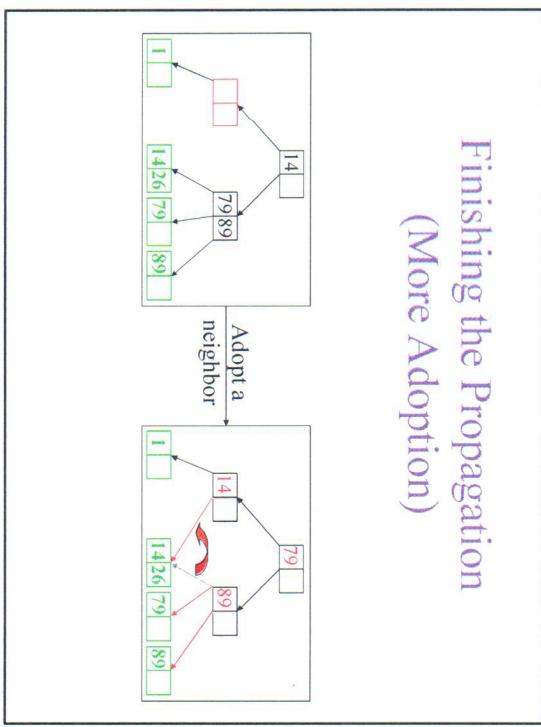
Deletion



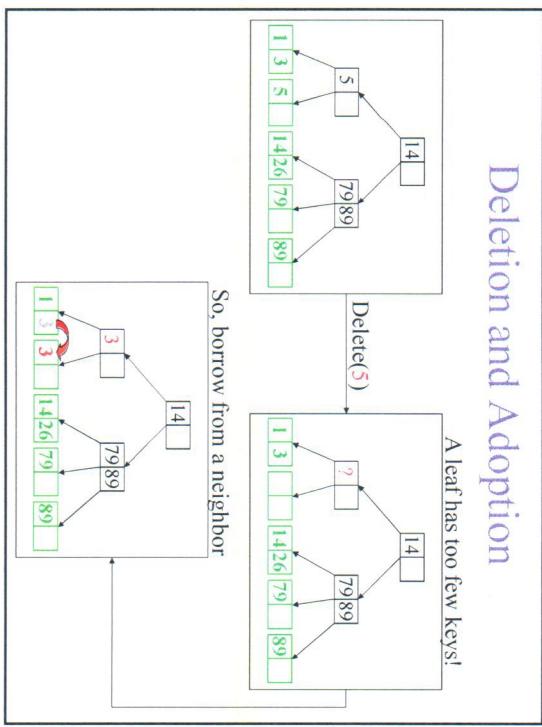
Deletion with Propagation



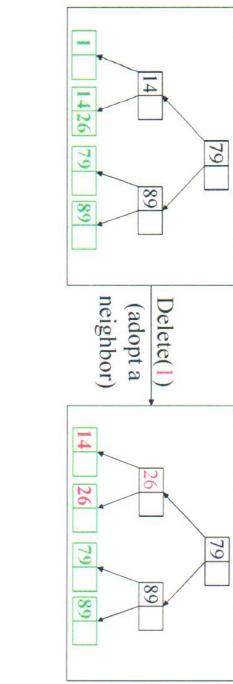
Finishing the Propagation (More Adoption)



Deletion and Adoption



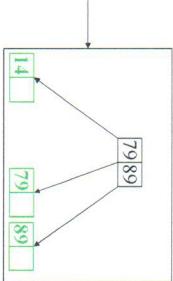
A Bit More Adoption



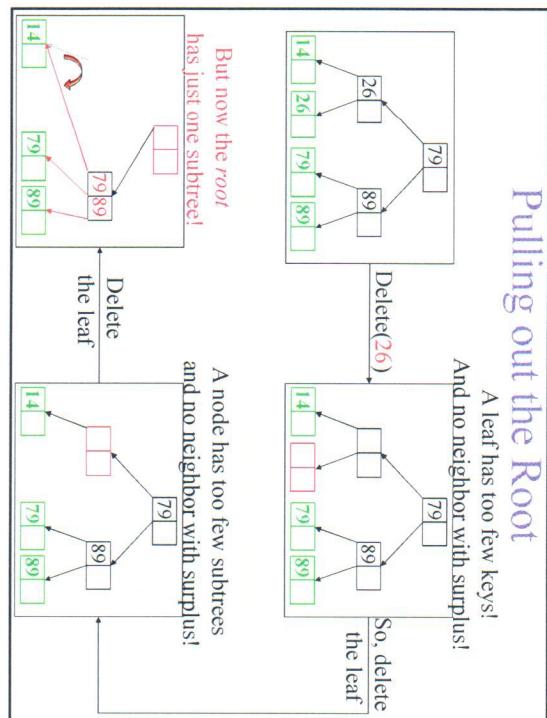
Pulling out the Root (continued)

The *root*
has just one subtree!

Just make
the one child
the new root!



But that's silly!



Deletion in Two Boring Slides of Text

- Remove the key from its leaf
 - If the leaf ends up with fewer than $\lceil \frac{m}{2} \rceil$ items, **underflow**!

- Adopt data from a neighbor:
 - If borrowing won't work, delete node and divide keys between neighbors

Why will dumping keys
always work if borrowing
doesn't?

- If the parent ends up with fewer than $\lceil \frac{m}{2} \rceil$ items, **underflow**!

Deletion Slide Two

- If a node ends up with fewer than $\lceil M/2 \rceil$ items, **underflow!**
 - Adopt subtrees from a neighbor; update the parent
 - If borrowing won't work, delete node and divide subtrees between neighbors
 - If the parent ends up with fewer than $\lceil M/2 \rceil$ items, **underflow!**
- If the root ends up with only one child, make the child the new root of the tree

This reduces the height of the tree!

Thinking about B-Trees

- B-Tree insertion can cause (expensive) splitting and propagation
- B-Tree deletion can cause (cheap) borrowing or (expensive) deletion and propagation
 - Propagation is rare if M and L are large (*Why?*)
- Repeated insertions and deletion can cause thrashing
- If $M = L = 128$, then a B-Tree of height 4 will store at least 30,000,000 items

A Tree with Any Other Name

FYI:

- B-Trees with $M = 3, L = 2$ are called 2-3 trees
- B-Trees with $M = 4, L = 3$ are called 2-3-4 trees

Why would we ever use these?

Another Way

B-Trees are *hairy* in practice

hairy - Slang: Fraught with difficulties; hazardous;
a hairy escape; hairy problems.

To Do

- Continue Project II
- Read chapter 4 in the book
- Look forward to no quiz/homework for two weeks!
- (And prepare for the midterm)

Coming Up

- Self-balancing Trees
 - AVL trees
 - Splay trees
- Second project due (January 31st)
- Midterm (February 4th)!