

Part II


Data Structures

§ Three Types of Data Structures Considered:

- Hash Tables
- Trees
- Heaps

• Hash Table


- suitable for large key space with small numbers of actual keys stored
- average search time over the hash table equal to $O(1)$
- collision-handling strategy required
- keys must be stored to ensure right items retrieved
- hash function maps a key to one table index



Hash Tables (continued)

- **Hash Table with Collision Resolution by Chaining**

- multiple elements with different keys mapped to the same table entry (collisions)
- they can be chained, with a lookup going through the chain
- for load factor of $\alpha = n/m$, where n (or m) is the number of elements (or table entries)
- under uniform distribution, one search (be successful or not) takes $\Theta(1+\alpha)$,
according to **Theorems 11.1 and 11.2**
- if hash table size (m) grows in proportion to n , then the mean search time is $\Theta(1)$
- there are other more efficient ways for handling collisions



Hash Tables (continued)

- **Hash Functions**

- division method
- multiplication method

- 1. Division method:** (more restrictive on the suitable m)

$h(k) = k \bmod m$, where the choice of $m = 2^p - 1$ is poor but
a prime not close to an exact power of 2 is often good

For example, given $n = 2000$ character strings to be stored in a hash table with
 $m = 701$ entries, then an unsuccessful search takes some 3 accesses

Hash Tables (continued)


2. Multiplication method:

- Multiply k by a constant A , with $0 < A < 1$ and extract the fraction part of $k \cdot A$, e.g.,
 $h(k) = \lfloor m (kA \bmod 1) \rfloor$, where " $kA \bmod 1$ " equals $kA - \lfloor kA \rfloor$

- table size (m) can be arbitrary, e.g., $m = 2^p$
- let A be a fraction of the form $s/2^w$ as shown in the figure below
- perform multiplication of w -bit k and w -bit s , to get a $2w$ -bit product
- product denoted by $r_1 \cdot 2^w + r_0$, then p -bit hash value is obtained from r_0
- while arbitrary A works, $A \approx (\sqrt{5} - 1)/2 = 0.6180339887 \dots$ recommended

Golden ratio: a root
of $x^2 + x - 1 = 0$.

+ Let $k = 123456$, $p = 14$, $m = 2^{14}$, $w = 32$
and $A = 2654435769/2^{32}$, we have
 $k \cdot s = (76300 \cdot 2^{32}) + 17612864$ to get $h(k) = 67$.



Hash Tables (continued)

- **Open Addressing** (to deal with collision-chaining)
 - elements stored inside the table
 - calculating probe sequence of given key, instead of using pointers:
 $\langle h(k, 0), h(k, 1), h(k, 2), \dots, h(k, m-1) \rangle$
 - if probe sequence is a permutation of $\langle 0, 1, 2, \dots, m-1 \rangle$, every table entry is a candidate location for the element
 - the probe sequence is fixed for a given key
 - key has to be stored in the table entry

HASH-INSERT(T, k)

```
 $i = 0$   
repeat  
     $j = h(k, i)$   
    if  $T[j] == \text{NIL}$   
         $T[j] = k$   
        return  $j$   
    else  $i = i + 1$   
until  $i == m$   
error “hash table overflow”
```

HASH-SEARCH(T, k)

```
 $i = 0$   
repeat  
     $j = h(k, i)$   
    if  $T[j] == k$   
        return  $j$   
     $i = i + 1$   
until  $T[j] == \text{NIL}$  or  $i = m$   
return NIL
```

Hash Tables (continued)

- **Uniform Hashing Analysis on Open Addressing**

- key probe sequence equal likely in one of $m!$ permutations of $\langle 0, 1, 2, \dots, m-1 \rangle$

- probe sequences defined for open addressing

- + linear probing: m different sequences

- dictated by $h(k, i) = (h'(k) + i) \bmod m$
- subject to *primary clustering*

- + quadratic probing: m different sequences

- dictated by $h(k, i) = (h'(k) + c_1 \cdot i + c_2 \cdot i^2) \bmod m$
- subject to *secondary clustering*, i.e., $h(k_1, 0) = h(k_2, 0) \Rightarrow$ both keys follow the same probe sequence

- + double hashing: m^2 different sequences

- dictated by $h(k, i) = (h_1(k) + i \cdot h_2(k)) \bmod m$
- many good choices: $h_2(k)$ relatively prime to m ; m itself a prime; m a power of 2 and $h_2(k)$ always an odd number; or below:

m a prime and m' slightly less than m (e.g., $m-1$ or $m-2$) with

$$h_1(k) = k \bmod m$$

$$h_2(k) = 1 + (k \bmod m')$$

e.g., for $k=14$ under $m=13$, $m'=11$;
upon inserting $k=14$.

0	
1	79
2	
3	
4	69
5	98
6	
7	72
8	
9	14
10	
11	50
12	

Hash Tables (continued)

• Open-Address Hashing Analysis

Theorem 1:

For an open-address hash table with load factor $\alpha = n/m < 1$, the expected number of probes in unsuccessful search under uniform hashing is at most $1/(1 - \alpha)$.

Note: unsuccessful search ends at seeing an empty entry.

- + if α is a constant, an unsuccessful search runs in $O(1)$ time
- + inserting an element into an open-address hash table with load factor α takes at most $1/(1 - \alpha)$ probes on an average under uniform hashing, because it has $(1-\alpha)$ for 1 probe, plus prob. $\alpha(1-\alpha)$ to take 2 probes, plus prob. $\alpha^2(1-\alpha)$ to take 3 probes, etc., yielding $1 + \alpha + \alpha^2 + \alpha^3 + \dots = 1/(1 - \alpha)$

e.g., for $\alpha = 50/100$, we have 2 probes.

Theorem 2:

For an open-address hash table with load factor $\alpha < 1$, the expected number of probes in a successful search under uniform hashing is at most $\frac{1}{\alpha} \cdot \ln \frac{1}{1 - \alpha}$.

- successful search for a key equals the sequence of inserting the key

2nd key insertion takes $\leq 1/(1-(1/m))$ probes, on an average, when α is $1/m$

3rd key takes $\leq 1/(1-(2/m))$ probes, on an average, when α is $2/m$

:

$i+1^{\text{th}}$ key takes $\leq 1/(1-(i/m))$ probes

Mean no. of probes equals results over all n keys inserted

This is because **insertion and probes** follow the same hash function.

We have: $1/n (\sum_{i=0}^{n-1} 1/(1 - i/m))$

e.g., for $\alpha = 50/100$, we have ≈ 1.39 probes, as $\ln(2) \sim 0.693$. (base $e \sim 2.718$)

Hash Tables (continued)

- **Universal Hashing**

For the hash function of $h_{ab}(k) = ((a \cdot k + b) \bmod p) \bmod m$, where p is a large prime number, with $p > m$, $a \in \{1, 2, \dots, p-1\}$ and $b \in \{0, 1, 2, \dots, p-1\}$, the collection of such hash functions is *universal*.

In this case, m can be any number and does not have to be a prime.
For example, for $p = 17$ and $m = 8$, we have $h_{34}(15) = 7$.

$$((3 \cdot 15 + 4) \bmod 17) \bmod 8 = 7$$

Hash Tables (continued)

● Perfect Hashing

- set of keys is *static* (never changed), then excellent worst case performance exists
- search takes $O(1)$ in memory accesses in the worst case
- two levels of hashing, with secondary hash table S_j to hold those hashed to slot j
- unlikely to have collisions in S_j (with n_j keys) if its size m_j equals $n_j^2 \rightarrow$ easy to determine a collision-free secondary hash function
- total storage required for both primary and secondary hash tables is $O(n)$

Theorem 3:

Given n_j keys for a hash table sized $m_j = n_j^2$, the probability of having collisions $< 1/2$.

This theorem implies that one can easily (with a few trials) identify **collision-free secondary** $h_j(k)$ with proper a_j and b_j

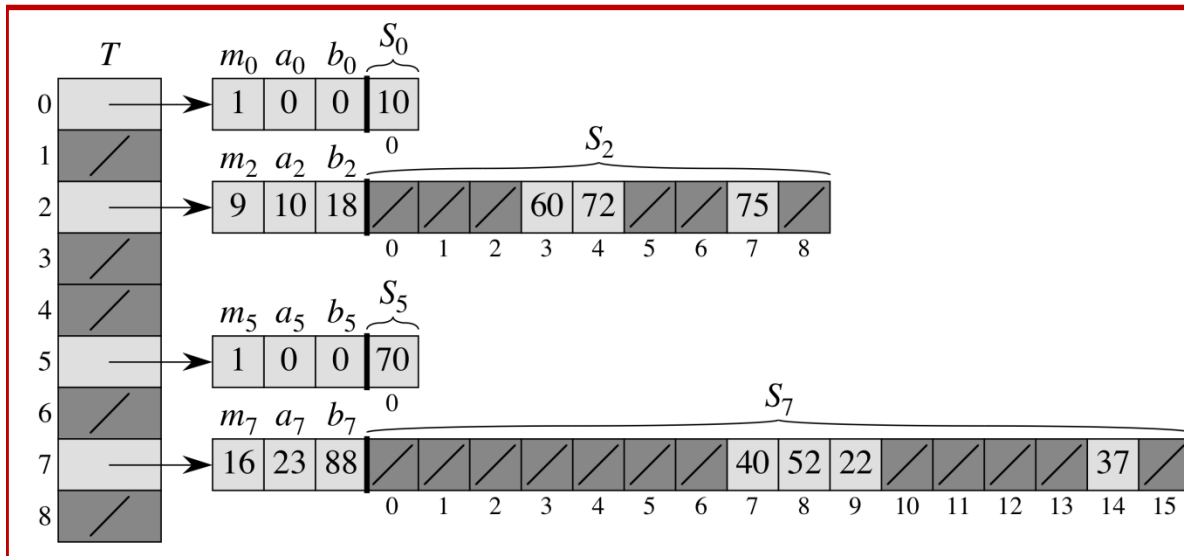


Figure 11.6 Using perfect hashing to store the set $K = \{10, 22, 37, 40, 60, 70, 75\}$. The primary hash function is $h(k) = ((a \cdot k + b) \bmod p) \bmod m$, where $a = 3$, $b = 42$, $p = 101$, and $m = 9$.

For example, $h(75) = 2$, so Key 75 hashes to slot 2 of table T . A secondary hash table S_j stores all keys hashing to Slot j . The size of hash table S_j is m_j , and the associated hash function is $h_j(k) = ((a_j k + b_j) \bmod p) \bmod m_j$. If $h_2(75) = 7$, Key 75 is stored in Slot 7 of secondary hash table S_2 . There are no collisions in any of the secondary hash tables, and so searching takes a constant time in the worst case.

Hash Tables (continued)

- **Perfect Hashing**

- total expected storage for both primary and secondary hash tables is $O(n)$.

Theorem 4:

Given n keys for an outer hash table sized $m = n$ and n_j keys present in Slot j under outer hash, we have $E \left[\sum_{j=0}^{m-1} n_j^2 \right] < 2n$. (Note: this deals with the "expected" storage size.)

Corollary 1:

With n keys for a hash table sized $m = n$ and n_j keys present in Slot j under the outer hash and the sizes of secondary hash tables equal to $m_j = n_j^2$, we have the probability of total storage used for secondary hash tables $\geq 4n$ less than $\frac{1}{2}$.

This corollary implies that one can easily (with a few trials) identify **good primary hash** $h(k)$ with proper a and b that lead to a reasonable amount of total storage required.

Hash Tables (continued)

- **Cuckoo Hashing**[§]


- move keys around upon insertion for worst probe case of $O(1)$
 - + quick search but possibly lengthy insertion
 - + multiple hash functions required for separate hash tables
- delete keys in the worst-case of $O(1)$

- example with two hash functions below:

- + two hash tables, each with m elements
- + two hash functions, h_1 and h_2 , one for a table
- + each key x at either $h_1(x)$ or $h_2(x)$
- + new key '10' with $h_1(10) = 6$ and $h_2(10) = 4$
 - * move key '93' around, if $h_2(93)=5$, for insertion

- Algorithm

1. insert "new key" to T_1 located at $h_1(\text{new key})$, when available
2. otherwise, "new key" displaces "existing key" in T_1 ; place "existing key" in T_2 located at $h_2(\text{existing key})$; repeat the displacement process, if needed.



§ R. Pagh and F. Rodler, "Cuckoo Hashing," *Journal of Algorithms*, Aug. 2001, pp. 121-133.

Hash Tables (continued)

- **Cuckoo Graph**

- derived from cuckoo hash tables

- + each table entry is a node

- + each key is an edge, which links the entries that can hold the key

- + an insertion adds a new edge to the graph

- * let $h_1(88) = 4$ and $h_2(88) = 6$

- * displace key '58' in T_1 to make room for key '88'

- * displace key '91' in T_2 to make room for key '58'


- * displace key '16' in T_1 to make room for key '91'

- * displace key '10' in T_2 to make room for key '16'

- * repeat

- + insertion done by tracing a path over the graph

- + what about inserting next key "67" with
 $h_1(67) = 1$ and $h_2(67) = 6$?




Hash Tables (continued)

- **Cuckoo Graph**

- insertion succeeds if and only if the new edge (defined by the new key) contains at most one cycle

Proof sketch:

1. each edge denotes a key and requires a table entry (i.e., node) to hold it
2. for a cycle exists, the number of its edges equals the number of its nodes involved, so that all nodes (table entries) are taken
3. if a new edge (after added) is on two cycles the edge must link two taken nodes AND all nodes on the two cycles have been used
4. the new key (edge) *cannot* be accommodated.



Side note: two cycles (one with 4 distinct nodes and the other with at least one distinct node) that share one edge include (at least) 5 distinct nodes but contain (at least) 6 distinct edges; impossible.

Binary Search Trees

§ Binary Search Tree Property

- Given tree node x , if node y is in the left (or right) subtree of x , we have $y.key \leq x.key$ (or $y.key \geq x.key$).

Theorem 1

INORDER-TREE-WALK(x) across an n -node subtree rooted at x takes $\Theta(n)$.

(because the tree height may equal n)

Proof.

The tree walk has to visit every node, and hence, the time complexity $T(n)$ is lower bounded by $\Omega(n)$.

Next, considering that INORDER-TREE-WALK(x) is called on x whose left and right subtrees have k and $(n-k-1)$ nodes, respectively, we have

$$T(n) \leq T(k) + T(n-k-1) + d \text{ for some } d > 0.$$

By the substitution method, we prove that $T(n) \leq (c+d)n + c$, for a constant c , below:

$$\begin{aligned} T(n) &\leq T(k) + T(n-k-1) + d \\ &\leq ((c+d)k + c) + ((c+d)(n-k-1) + c) + d \\ &= (c+d)n + c - (c+d) + c + d \\ &= \underline{(c+d)n + c} \quad (\text{upper bounded}) \end{aligned}$$

Binary Search Trees (continued)

- **Querying binary search trees**

- searching
- successor and predecessor
- insertion and deletion

– Searching via TREE-SEARCH in time complexity $O(h)$

TREE-SEARCH(x, k)


if $x == \text{NIL}$ or $k == \text{key}[x]$

return x

if $k < x.\text{key}$

return TREE-SEARCH($x.\text{left}, k$)

else return TREE-SEARCH($x.\text{right}, k$)



search path from root to $\text{key} = 13$ is
 $15 \rightarrow 6 \rightarrow 7 \rightarrow 13$


Binary Search Trees (continued)

- Successor and predecessor
 - + by inorder tree walk
 - + separate results for non-empty right subtree and for empty right subtree
 - + time complexity for a tree of height h is $O(h)$

TREE-SUCCESSOR(x)

```
if  $x.right \neq \text{NIL}$ 
    return TREE-MINIMUM( $x.right$ )
 $y = x.p$     y is the parent node of x.
while  $y \neq \text{NIL}$  and  $x == y.right$ 
     $x = y$     climb up the tree.
     $y = y.p$ 
return  $y$ 
```

If node x has no right subtree, its successor is the lowest ancestor of x whose left subtree contains x .



successor of $key = 7$ is 9
successor of $key = 13$ is 15

Binary Search Trees (continued)

– Insertion

- + tree node (say, z) to be added, with $z.key = v$, $z.left = \text{nil}$, $z.right = \text{nil}$
- + z added to one appropriate node with an **absent** left or right child (never a full node)
- + tree in-order walk, with the trailing pointer y maintained
- + time complexity for a tree of height h is $\underline{O(h)}$

TREE-INSERT(T, z)

$y = \text{NIL}$

$x = T.\text{root}$

while $x \neq \text{NIL}$

$y = x$

if $z.key < x.key$

$x = x.left$

else $x = x.right$

y is the parent node of x after
 x is reassigned to move down

tree traversal downward until
the insertion point, with the
added node as a **leaf node**

$z.p = y$ // y is the parent node of z .


if $y == \text{NIL}$

$T.\text{root} = z$ // tree T was empty

elseif $z.key < y.key$

$y.left = z$

else $y.right = z$



Inserting $\text{key} = 13$ into binary search tree, with affected path marked.

Binary Search Trees (continued)

– Deletion (on z)

- + three basic cases involved: z having no child, exactly one child, two children
- + based on TRANSPLANT that replaces sibling subtree with another subtree
- + time complexity for a tree of height h is $O(h)$

TRANSPLANT(T, u, v)

```

if  $u.p == \text{NIL}$ 
     $T.\text{root} = v$ 
elseif  $u == u.p.\text{left}$ 
     $u.p.\text{left} = v$ 
else  $u.p.\text{right} = v$ 
if  $v \neq \text{NIL}$ 
     $v.p = u.p$ 
    
```


TREE-DELETE(T, z)

```

if  $z.\text{left} == \text{NIL}$ 
    TRANSPLANT( $T, z, z.\text{right}$ )    //  $z$  has no left child
elseif  $z.\text{right} == \text{NIL}$ 
    TRANSPLANT( $T, z, z.\text{left}$ )     //  $z$  has just a left child
else //  $z$  has two children.
     $y = \text{TREE-MINIMUM}(z.\text{right})$  //  $y$  is  $z$ 's successor
    if  $y.p \neq z$  ← This is case (d) shown in the figure.
        //  $y$  lies within  $z$ 's right subtree but is not the root of this subtree
        TRANSPLANT( $T, y, y.\text{right}$ )
         $y.\text{right} = z.\text{right}$ 
         $y.\text{right}.p = y$ 
        // Replace  $z$  by  $y$ .
        TRANSPLANT( $T, z, y$ )
         $y.\text{left} = z.\text{left}$ 
         $y.\text{left}.p = y$ 
    
```

Prepare a subtree rooted at y whose left subtree is empty.

Alternatively, one may cut the left-hand subtree of z and move it to become the left-hand subtree of y .



Binary Search Trees (continued)

- Randomly build binary search trees
 - + worst-case tree height (h) being $n-1$
 - + can be shown that $h \geq \lfloor \lg n \rfloor$, which is the best case
 - + like quicksort, average case proven to be much closer to best case than worst case
 - + special case of creating binary search trees (via insertion alone) with random keys, we have following theorem:

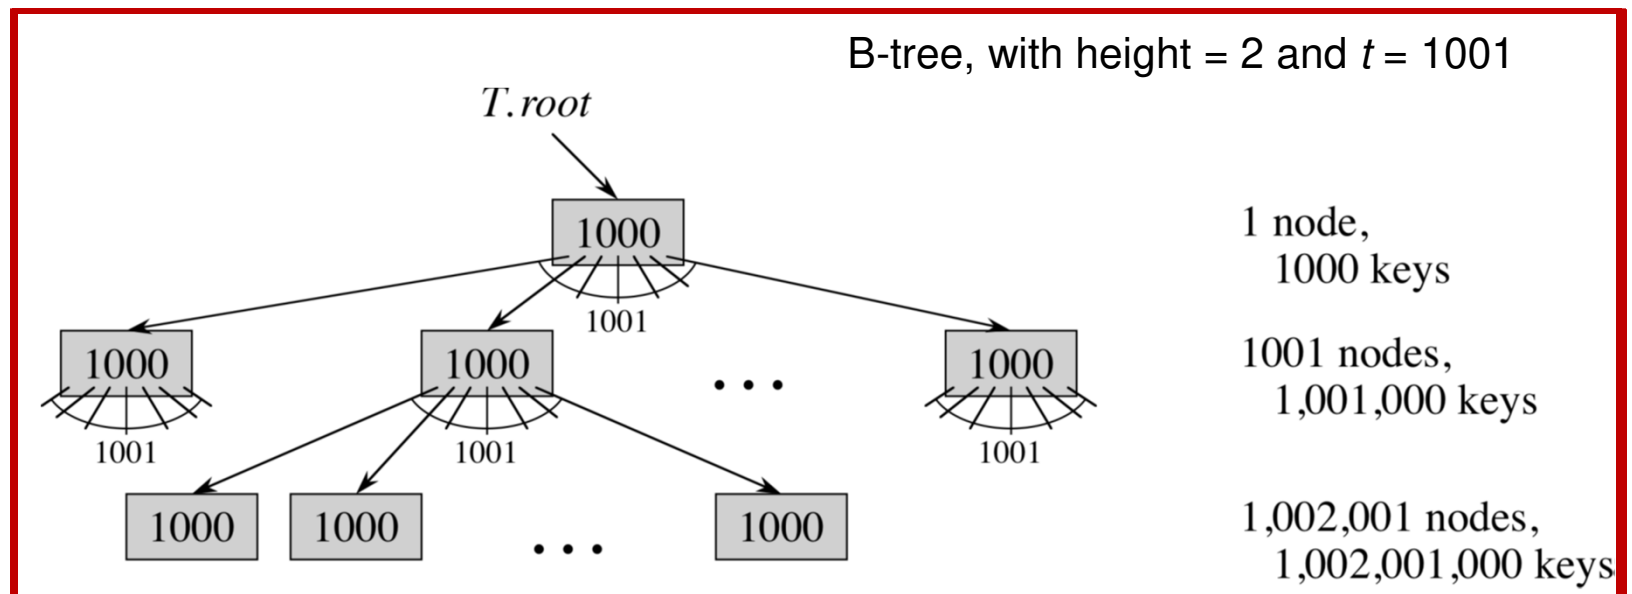
Theorem 2.

← This theorem refers to the special case that the tree is created by insertion alone (and no deletion).

The expected height of a randomly built binary search tree on n distinct keys is **$O(\lg n)$** .

B-Trees

- § **Balanced Search Trees (B-Trees)** with t — non-root node has $\geq t$ & $\leq 2t$ children
- + balance achieved by keeping multiple (i.e., $\geq t-1$) keys in each non-root node
 - + node has at most $2t-1$ keys, called **full node** (with degree = $2t$)
 - + keys stored in a node in non-decreasing order
 - + node x (with $x.n$ keys) has $x.n+1$ children, pointed by $x.c_1, x.c_2, \dots, x.c_{x.n+1}$, then $k_1 \leq x.key_1 \leq k_2 \leq x.key_2 \leq \dots \leq x.key_{x.n} \leq k_{x.n+1}$, where k_i is a key stored in subtree rooted at $x.c_i$ and $x.key_i$ is a key stored in node x
 - + all leaves have the same height
 - + simplest B-tree is for $t = 2$, called 2-3-4 tree (each *internal node* has 2, 3, or 4 children)

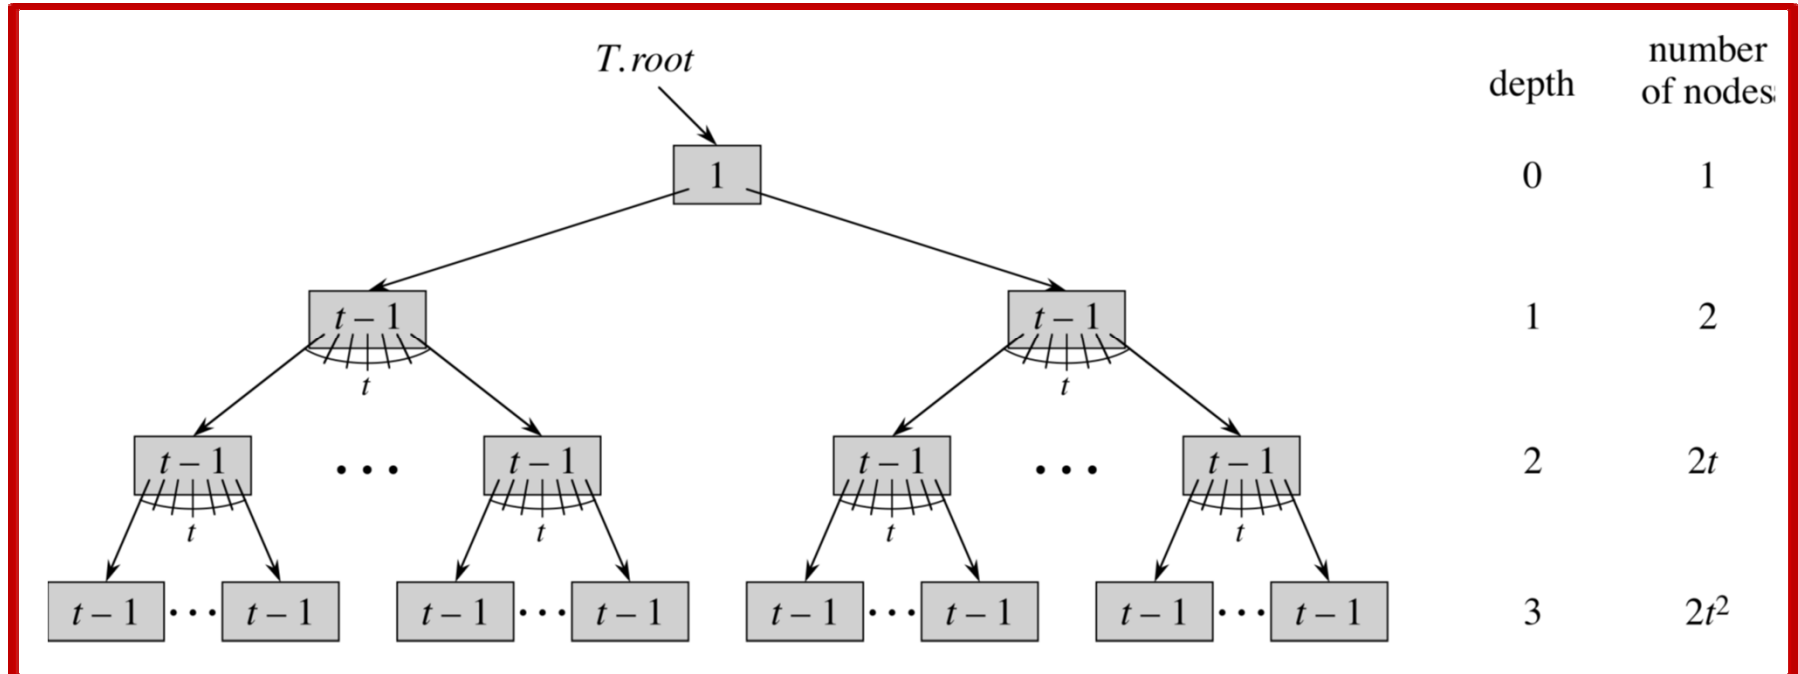


B-Trees (continued)

• Theorem

For any n -key B-tree T of height h and minimum node degree $t \geq 2$ (i.e., number of keys in each non-root node $\geq t-1$), we have $h \leq \log_t \frac{n+1}{2}$.

+ Proof follows the fact of $n \geq 1 + (t-1) \cdot \sum_{i=1}^h 2 \cdot t^{i-1}$ illustrated in the figure below (to show the minimum possible number of keys in a B-tree with height = 3).



B-Trees (continued)

- **Basic operations**

- + Searching for record with key = k : each internal node x makes an $(x.n + 1)$ -way branching decision; returning $(y, i) \rightarrow$ node y and its index i with $y.key_i = k$

B-TREE-SEARCH(x, k)

```
1   $i = 1$ 
2  while  $i \leq x.n$  and  $k > x.key_i$ 
3       $i = i + 1$ 
4  if  $i \leq x.n$  and  $k == x.key_i$ 
5      return  $(x, i)$ 
6  elseif  $x.leaf$ 
7      return NIL
8  else DISK-READ( $x.c_i$ )
9      return B-TREE-SEARCH( $x.c_i, k$ )
```

$x.leaf$: Boolean variable to indicate if it is a leaf node

B-Trees (continued)


• Basic operations

- + **Splitting a full node**: for *non-full internal node* x and a *full child* of x (say, $y = x.c_i$), y is split at median key S , which moves up into x , with those $> S$ forming a new subtree

B-TREE-SPLIT-CHILD(x, i) (split i^{th} child of x)

```

1   $z = \text{ALLOCATE-NODE}()$ 
2   $y = x.c_i$  (Node  $x$ 's  $i^{\text{th}}$  child is full)
3   $z.\text{leaf} = y.\text{leaf}$  (status – leaf?)
4   $z.n = t - 1$ 
5  for  $j = 1$  to  $t - 1$ 
6       $z.\text{key}_j = y.\text{key}_{j+t}$ 
7  if not  $y.\text{leaf}$ 
8      for  $j = 1$  to  $t$  (move child pointers from  $y$  to  $z$ 
9           $z.c_j = y.c_{j+t}$  with pointer # from 1 to  $t$ )
10  $y.n = t - 1$ 
11 for  $j = x.n + 1$  downto  $i + 1$ 
12      $x.c_{j+1} = x.c_j$  (shift child pointers in  $x$  rightward)
13  $x.c_{i+1} = z$ 
14 for  $j = x.n$  downto  $i$ 
15      $x.\text{key}_{j+1} = x.\text{key}_j$  (shift keys in  $x$  rightward)
16  $x.\text{key}_i = y.\text{key}_t$  (move middle key in  $y$  up)
17  $x.n = x.n + 1$ 
18 DISK-WRITE( $y$ )
19 DISK-WRITE( $z$ )
20 DISK-WRITE( $x$ )
    
```



B-Trees (continued)

• Basic operations

- + **Inserting a key**: B-tree T of height h , takes $O(h)$ disk accesses (nodes kept as disk pages)
 - *can't* insert the key into a newly created node *nor* into an internal node directly
 - so can insert it only to a **leaf node**
 - **never** descend through a **full node**, achieved by B-TREE-SPLIT-CHILD, to avoid back-tracking altogether
 - root split to increase height by 1

B-TREE-INSERT(T, k)

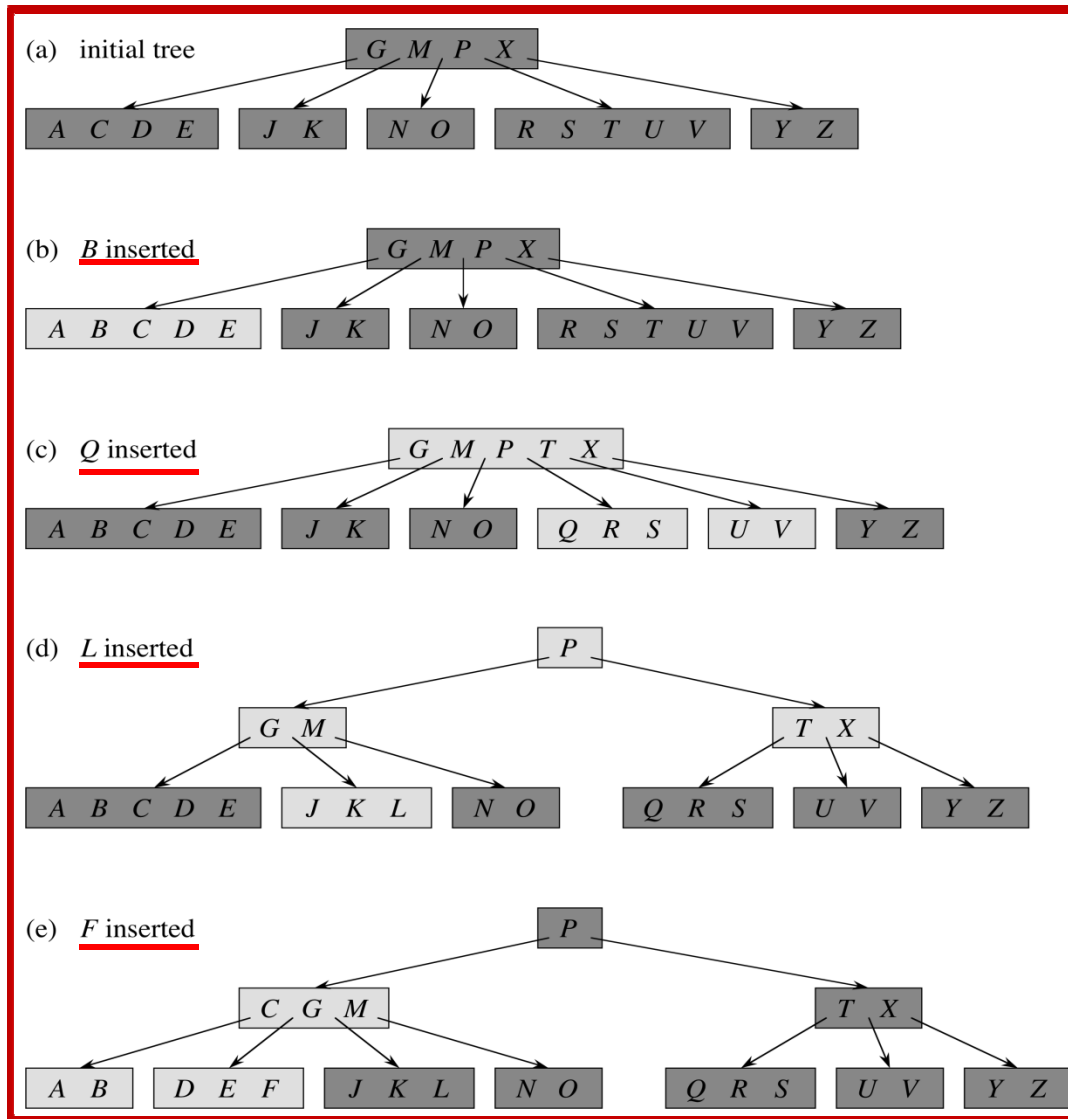
```
1   $r = T.root$ 
2  if  $r.n == 2t - 1$     // splitting the root?
3       $s = \text{ALLOCATE-NODE}()$ 
4       $T.root = s$ 
5       $s.leaf = \text{FALSE}$ 
6       $s.n = 0$ 
7       $s.c_1 = r$ 
8      B-TREE-SPLIT-CHILD( $s, 1$ ) (split 1st child of  $S$ )
9      B-TREE-INSERT-NONFULL( $s, k$ )
10 else B-TREE-INSERT-NONFULL( $r, k$ )
```

B-TREE-INSERT-NONFULL(x, k)

```
1   $i = x.n$ 
2  if  $x.leaf$  (upon a leaf, key is inserted therein)
3      while  $i \geq 1$  and  $k < x.key_i$ 
4           $x.key_{i+1} = x.key_i$  (shift key in  $x$  rightward)
5           $i = i - 1$ 
6           $x.key_{i+1} = k$ 
7           $x.n = x.n + 1$ 
8          DISK-WRITE( $x$ )
9  else while  $i \geq 1$  and  $k < x.key_i$ 
10      $i = i - 1$ 
11      $i = i + 1$ 
12     DISK-READ( $x, c_i$ ) (descend one level)
13     if  $x.c_i.n == 2t - 1$  // full child?
14         B-TREE-SPLIT-CHILD( $x, i$ )
15         if  $k > x.key_i$ 
16              $i = i + 1$ 
17     B-TREE-INSERT-NONFULL( $x.c_i, k$ )
```

B-Trees (continued)


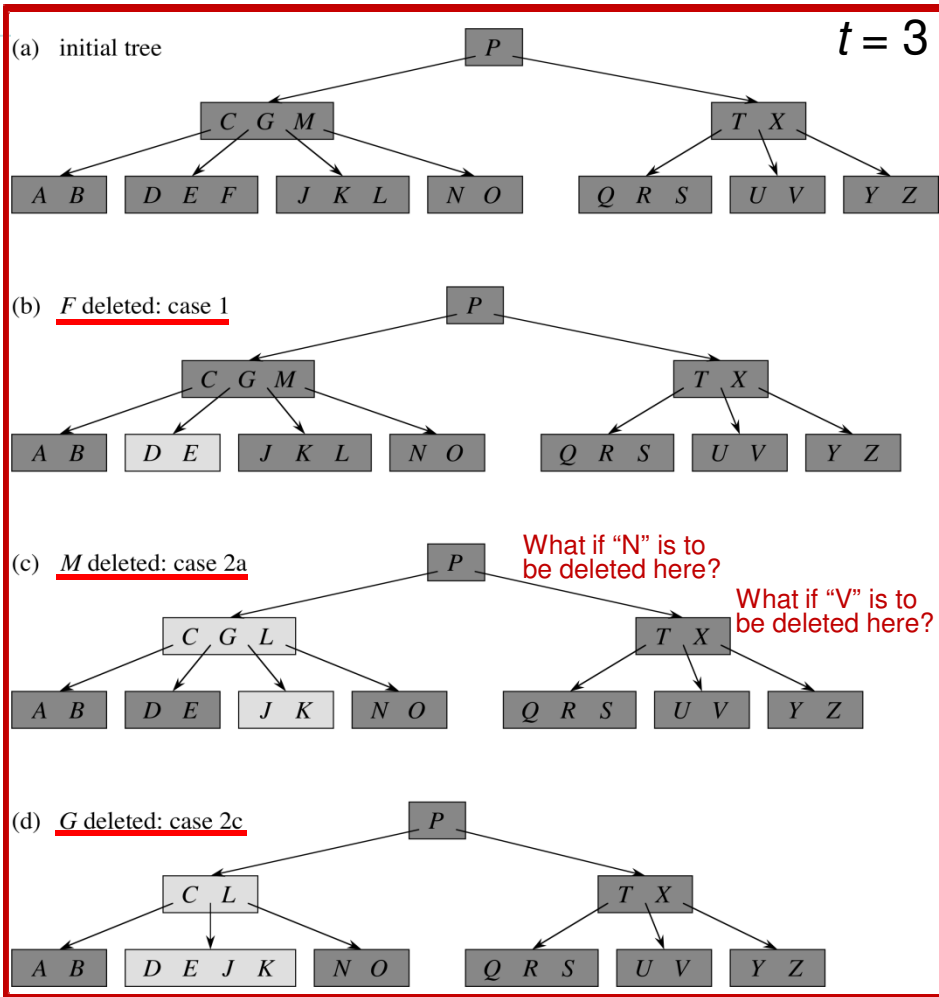
+ Inserting a key: B-tree T with $t = 3$ and modified nodes lightly shaded



B-Trees (continued)

• Basic operations

- + **Deleting a key**: delete key k from the subtree rooted at x in B-tree T with the minimum degree of t (where the key may be in any node, not just at a leaf node)
 - number of keys kept in any **internal node** upon descending through would be **at least t**



B-Trees (continued)

+ Deleting key k from the subtree rooted at x in B-tree T with the minimum degree of t
 – **Deletion Procedure** sketched below:

1. If the key k is in node x and x is a leaf, delete the key k from x .
2. If the key k is in node x and x is an internal node, do the following:
 - a. If the child y that precedes k in node x has at least t keys, then find the predecessor k' of k in the subtree rooted at y . Recursively delete k' , and replace k by k' in x . (We can find k' and delete it in a single downward pass.)
 - b. If y has fewer than t keys, then, symmetrically, examine the child z that follows k in node x . If z has at least t keys, then find the successor k' of k in the subtree rooted at z . Recursively delete k' , and replace k by k' in x . (We can find k' and delete it in a single downward pass.)
 - c. Otherwise, if both y and z have only $t - 1$ keys, merge k and all of z into y , so that x loses both k and the pointer to z , and y now contains $2t - 1$ keys. Then free z and recursively delete k from y .
3. If the key k is not present in internal node x , determine the root $x.c_i$ of the appropriate subtree that must contain k , if k is in the tree at all. If $x.c_i$ has only $t - 1$ keys, execute step 3a or 3b as necessary to guarantee that we descend to a node containing at least t keys. Then finish by recursing on the appropriate child of x .
 - a. If $x.c_i$ has only $t - 1$ keys but has an immediate sibling with at least t keys, give $x.c_i$ an extra key by moving a key from x down into $x.c_i$, moving a key from $x.c_i$'s immediate left or right sibling up into x , and moving the appropriate child pointer from the sibling into $x.c_i$.
 - b. If $x.c_i$ and both of $x.c_i$'s immediate siblings have $t - 1$ keys, merge $x.c_i$ with one sibling, which involves moving a key from x down into the new merged node to become the median key for that node.

Targeted node with the key

This is Rule 1 - Key k present:
 At internal node x which contains the key (say, k) to be deleted, it checks if predecessor or successor of k can be borrowed to replace k (and in this case, only the borrowed key is moved; and its associated pointer stays).

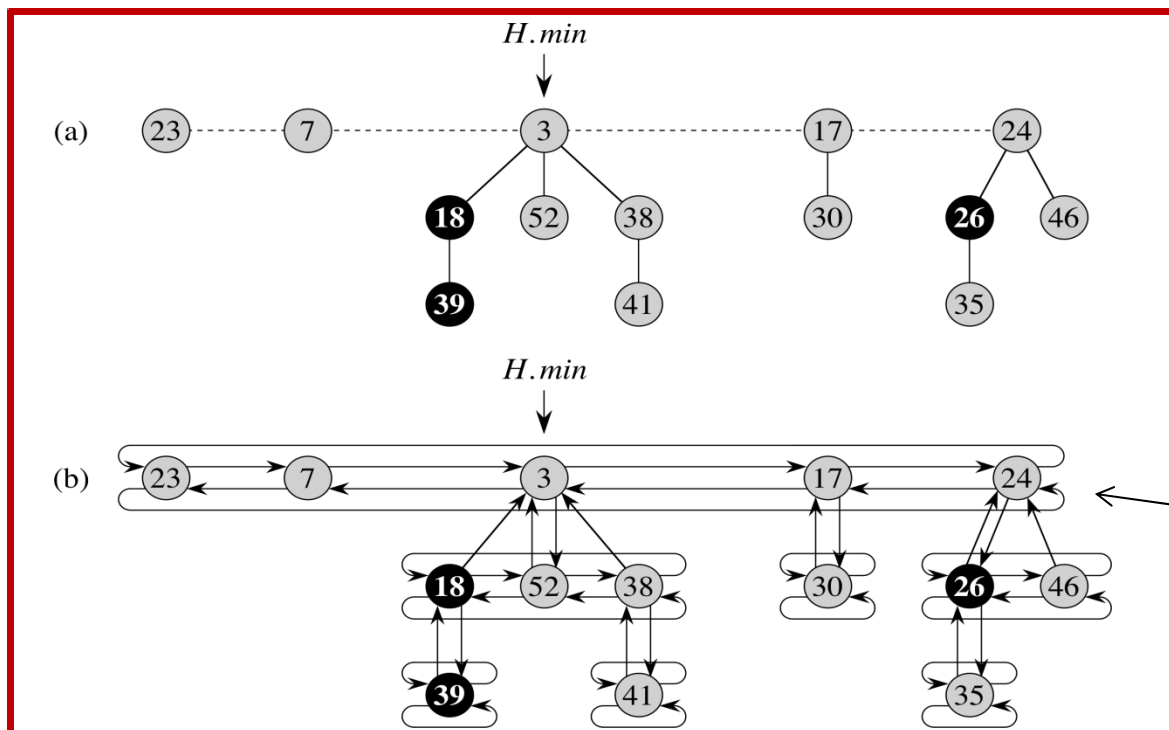
Prepare to walk down the tree

This is Rule 2 - Key k not present:
 While at internal node x , it checks if the root of the target child has at least t key
 If not, it prepares that root node by borrowing a key (plus associated pointer, i.e., a subtree) from its left or right sibling, if possible.

Fibonacci Heaps


§ Fibonacci Heap

- + a collection of min-heap trees, with their roots doubly linked together
- + belongs to the “mergeable heap” that supports five operations efficiently:
MAKE-HEAP(); INSERT(H, x); MINIMUM(H); EXTRACT-MIN(H); UNION(H_1, H_2)
- + additionally, the Fibonacci heap supports:
DECREASE-KEY(H, x, k); DELETE(H, x)
- + several preceding operations run in **constant** amortized time



Four pointers exist for each node:
 p , $child$, $left$, $right$, permitting
children of a node to be doubly
linked as the child list.

Max Binary Heap



Procedure	<u>Binary heap</u> (worst-case)	<u>Fibonacci heap</u> (amortized)
MAKE-HEAP	$\Theta(1)$	$\Theta(1)$
✦ INSERT (to Root List)	$\Theta(\lg n)$	$\Theta(1)$
MINIMUM	$\Theta(1)$	$\Theta(1)$
EXTRACT-MIN	$\Theta(\lg n)$	$O(\lg n)$
✦ UNION (i.e., Merging)	$\Theta(n)$	$\Theta(1)$
✦ DECREASE-KEY	$\Theta(\lg n)$	$\Theta(1)$
DELETE	$\Theta(\lg n)$	$O(\lg n)$

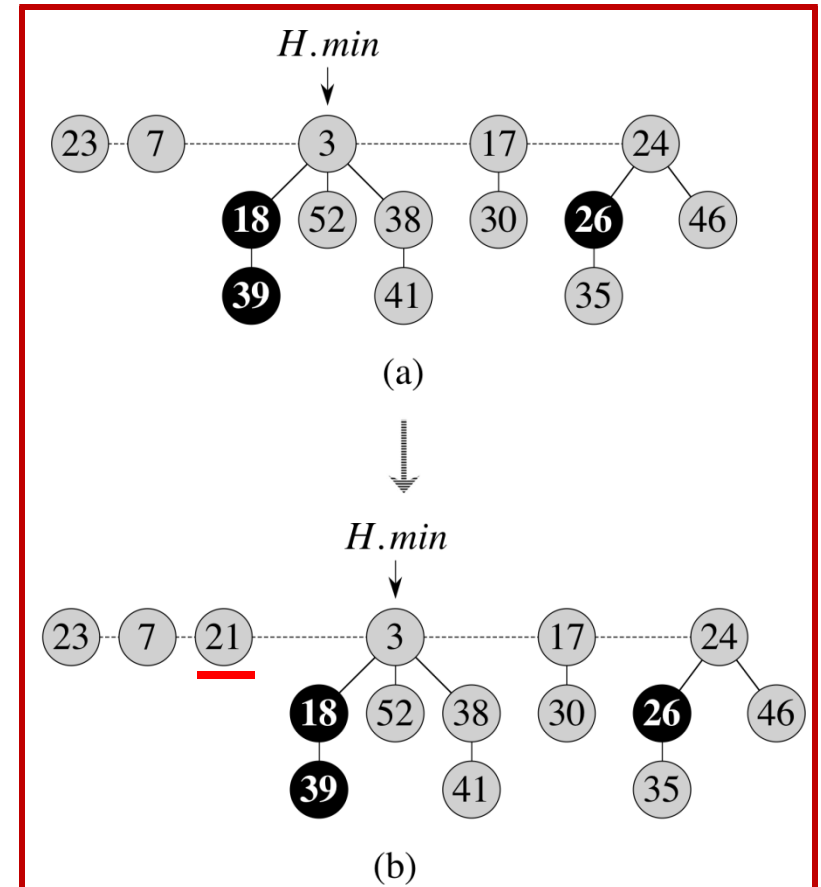
Fibonacci Heaps (continued)

- **Heap operations**

+ inserting a node, with time complexity $O(1)$, by adding the node to the root list

FIB-HEAP-INSERT(H, x)

```
1   $x.degree = 0$ 
2   $x.p = \text{NIL}$ 
3   $x.child = \text{NIL}$ 
4   $x.mark = \text{FALSE}$ 
5  if  $H.min == \text{NIL}$ 
6      create a root list for  $H$  containing just  $x$ 
7       $H.min = x$ 
8  else insert  $x$  into  $H$ 's root list
9      if  $x.key < H.min.key$ 
10          $H.min = x$ 
11   $H.n = H.n + 1$ 
```



Fibonacci Heaps (continued)

- **Uniting two Fibonacci heaps**

- + concatenating the root lists of H_1 and H_2
- + determining the new minimum node

FIB-HEAP-UNION(H_1, H_2)

```
1   $H = \text{MAKE-FIB-HEAP}()$ 
2   $H.min = H_1.min$ 
3  concatenate the root list of  $H_2$  with the root list of  $H$ 
4  if ( $H_1.min == \text{NIL}$ ) or ( $H_2.min \neq \text{NIL}$  and  $H_2.min.key < H_1.min.key$ )
5       $H.min = H_2.min$ 
6   $H.n = H_1.n + H_2.n$ 
7  return  $H$ 
```


Fibonacci Heaps (continued)

- **Extracting minimum node**

- + doing consolidation of trees on the root list, so that one root is left for each degree
- + calling CONSOLIDATE(H) to link trees of the same degree

conducting delayed work of consolidating trees in the root list

FIB-HEAP-EXTRACT-MIN(H)

```
1   $z = H.min$ 
2  if  $z \neq \text{NIL}$ 
3      for each child  $x$  of  $z$ 
4          add  $x$  to the root list of  $H$ 
5           $x.p = \text{NIL}$ 
6      remove  $z$  from the root list of  $H$ 
7      if  $z == z.right$            //  $z$  being the root of the only tree and having no child (since any child, if existing,
8           $H.min = \text{NIL}$            would then have been added to the root list)
9      else  $H.min = z.right$ 
10     CONSOLIDATE( $H$ )
11      $H.n = H.n - 1$ 
12 return  $z$ 
```

F_4

Fibonacci sequence: 0, 1, 1, 2, 3, 5, 8, 13, 21, 34, 55, ...

Fibonacci Heaps (continued)

- **CONSOLIDATE(H)** for consolidating trees on the root list

CONSOLIDATE(H)

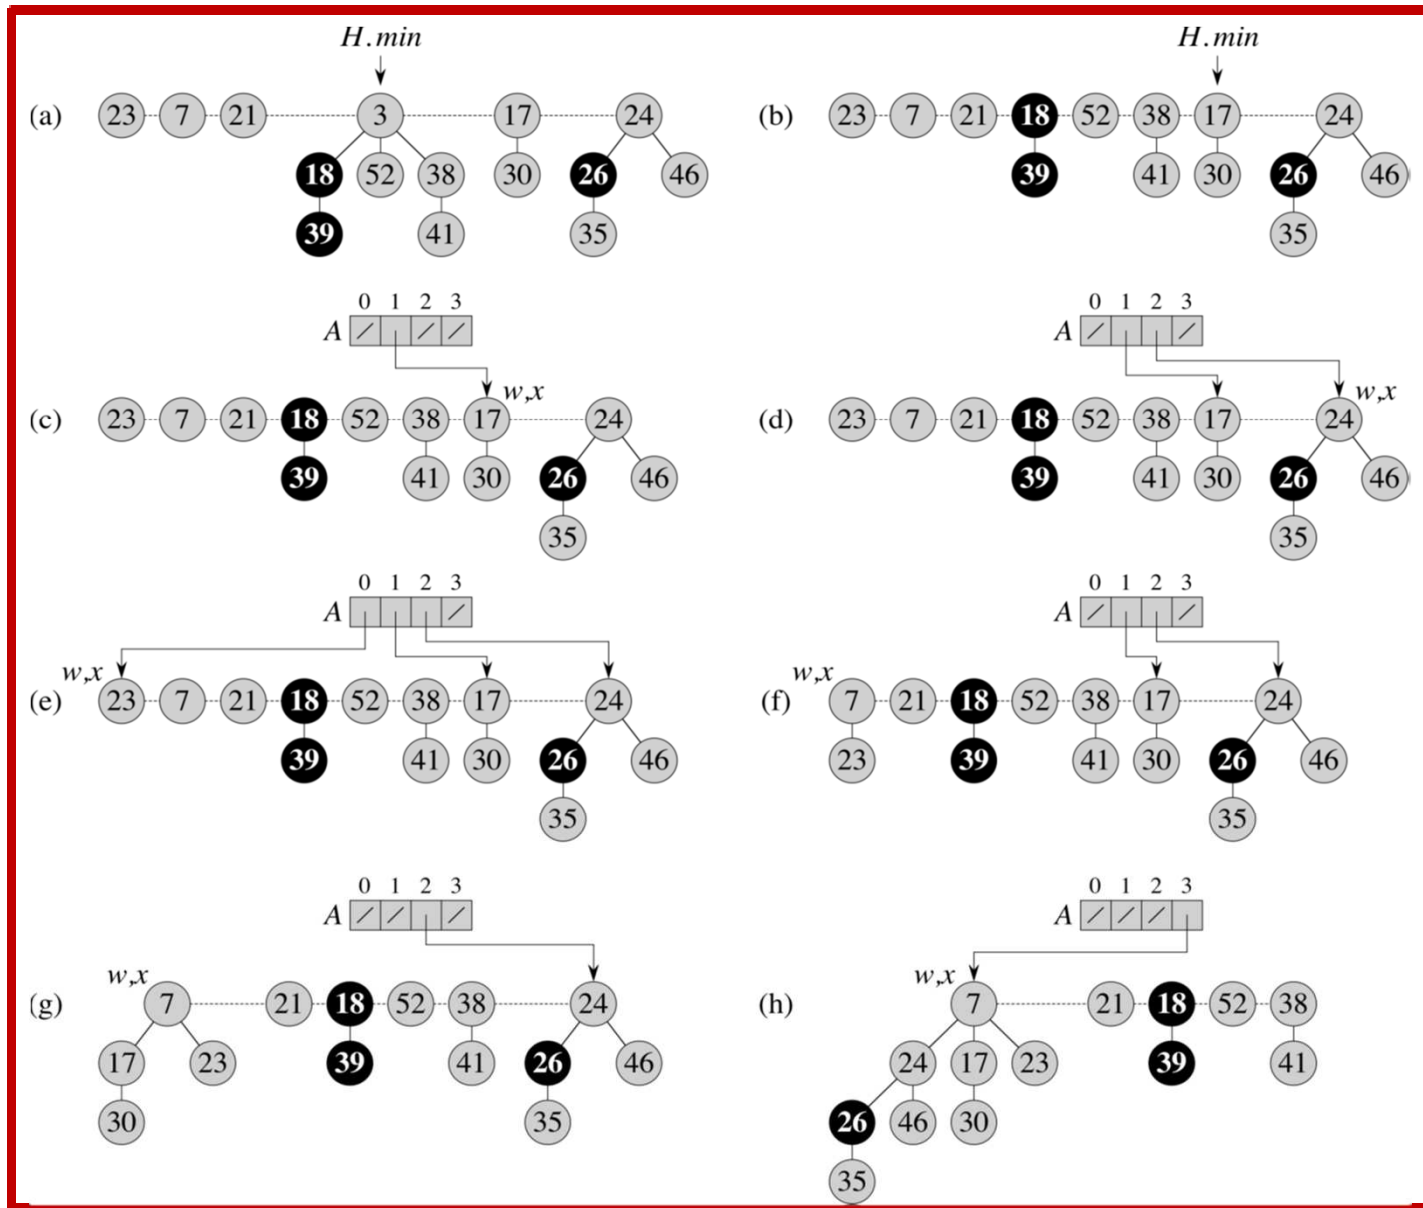
```
1  let  $A[0..D(H.n)]$  be a new array //  $D(H, n)$  being the largest degree of a constituent MIN heap
2  for  $i = 0$  to  $D(H.n)$ 
3       $A[i] = \text{NIL}$  // auxiliary array initialization
4  for each node  $w$  in the root list of  $H$ 
5       $x = w$ 
6       $d = x.\text{degree}$ 
7      while  $A[d] \neq \text{NIL}$ 
8           $y = A[d]$  // another node with the same degree as  $x$  existing
9          if  $x.\text{key} > y.\text{key}$ 
10             exchange  $x$  with  $y$ 
11             FIB-HEAP-LINK( $H, y, x$ )
12              $A[d] = \text{NIL}$ 
13              $d = d + 1$ 
14              $A[d] = x$ 
15   $H.\text{min} = \text{NIL}$ 
16  for  $i = 0$  to  $D(H.n)$  // establish the root list and find  $H.\text{min}$  after consolidation
17      if  $A[i] \neq \text{NIL}$ 
18          if  $H.\text{min} == \text{NIL}$ 
19             create a root list for  $H$  containing just  $A[i]$ 
20              $H.\text{min} = A[i]$ 
21          else insert  $A[i]$  into  $H$ 's root list
22              if  $A[i].\text{key} < H.\text{min}.\text{key}$ 
23                   $H.\text{min} = A[i]$ 
```

FIB-HEAP-LINK(H, y, x)

```
1  remove  $y$  from the root list of  $H$ 
2  make  $y$  a child of  $x$ , incrementing  $x.\text{degree}$ 
3   $y.\text{mark} = \text{FALSE}$ 
```






Fibonacci Heaps (continued)

- Extracting minimum node



Fibonacci Heaps (continued)

- Extracting minimum node



Fibonacci Heaps (continued)

- Decreasing a key

FIB-HEAP-DECREASE-KEY(H, x, k)

```
1  if  $k > x.key$ 
2      error "new key is greater than current key"
3   $x.key = k$ 
4   $y = x.p$ 
5  if  $y \neq \text{NIL}$  and  $x.key < y.key$ 
6      CUT( $H, x, y$ )
7      CASCADING-CUT( $H, y$ )
8  if  $x.key < H.min.key$ 
9       $H.min = x$ 
```

CUT(H, x, y)


```
1  remove  $x$  from the child list of  $y$ , decrementing  $y.degree$ 
2  add  $x$  to the root list of  $H$ 
3   $x.p = \text{NIL}$ 
4   $x.mark = \text{FALSE}$ 
```

CASCADING-CUT(H, y)

```
1   $z = y.p$ 
2  if  $z \neq \text{NIL}$ 
3      if  $y.mark == \text{FALSE}$ 
4           $y.mark = \text{TRUE}$ 
5      else CUT( $H, y, z$ )
6      CASCADING-CUT( $H, z$ )
```

Fibonacci Heaps (continued)

- Decreasing a key



Fibonacci Heaps (continued)

- **Deleting a key**

FIB-HEAP-DELETE(H, x)

FIB-HEAP-DECREASE-KEY($H, x, -\infty$)

FIB-HEAP-EXTRACT-MIN(H)

Fibonacci Heaps (continued)

• Heap Bounds

Lemma 1.

Let x be any node in a Fibonacci heap and $k = x.degree$. We have $size(x) \geq F_{k+2} \geq \phi^k$, where $\phi = (1 + \sqrt{5})/2$.

$= 1.618$

Fibonacci sequence: 0, 1, 1, 2, 3, 5, 8, 13, 21, 34, 55, ...
 ϕ^k 1.0, 1.6, 2.6, 4.2, 6.9, 11.1, 17.9, 29.0, 47.0


$k=1, k=2, k=3, k=4,$

In fact, $D(k) \geq$

2, 4, 8, 16, 32, 64, 128,

Lemma 2.

The maximum degree $D(x)$ of any node in an n -node Fibonacci heap is $O(\lg n)$.



Note: the smallest Fibonacci heaps with $k = 1, 2, 3, 4,$ etc. include 2, 4, 8, 16, etc. nodes. Hence, $size(1) \geq 2$, $size(2) \geq 4$, $size(3) \geq 8$, $size(4) \geq 16$, etc.